

# Assignment 2: Sentiment Classification Using Logistic Regression

## Programming Assignment (100 Points scaled to 40)

For this assignment we will be implementing a naive bayes baseline classifier. Additionally, we will be using pytorch to implement a binary logistic regression classifier. Our task is sentiment classification for hotel reviews. The input to your model will be a text review, and the output is a 1 for 0 marking it as positive or negative.

We have provided a utility file for loading the data, and some of the basic modeling. Your task is to fill in the functions below in order to train an accurate a classifier as possible!

We suggest browsing the utility script first. Additionally, make sure to install dependencies from the provided requirements.txt file in a similar fashion to the pytorch tutorial. With your environment activated in the terminal, run:

```
pip install -r requirements.txt
```

```
In [30]: from typing import List
import random
```

## Section 1: Sentiment Classification Dataset (Total: 20 Points)

The training data for this task consists of a collection of short hotel reviews. The data is formatted as one review per line. Each line starts with a unique identifier for the review (as in ID-2001) followed by tab and the text of the review. The reviews are not tokenized or sentence segmented in any way (the words are space separated). The positive reviews and negative reviews appear in separate files namely `hotelPos1-train.txt` and `hotelNeg1-train.txt`.

```
In [37]: from util import load_train_data
pos_datapath = "data/hotelPos1-train.txt"
neg_datapath = "data/hotelNeg1-train.txt"
all_texts, all_labels = load_train_data(pos_datapath, neg_datapath)
```

### Lets look at what is in the data

```
In [40]: def random_sample(texts, labels, label):
    data_by_label = {}
    for lab,text in zip(labels, texts):
        if lab not in data_by_label:
            data_by_label[lab] = []
        data_by_label[lab].append(text)
    return random.choice(data_by_label[label])

print("""--- Positive Example ---""")
print(random_sample(all_texts, all_labels, label=1))
print("""--- Negative Example ---""")
print(random_sample(all_texts, all_labels, label=0))
```

--- Positive Example ---  
I have no complaints and I'm a picky hotel guest. The room was clean, the sheets were not old, and the bathroom was spotless. Everything worked properly which made for a great experience while getting ready for an interview. The front desk staff was more than helpful. They gave us advice on where to eat, information on the city, and the place I was interviewing at. The continental breakfast had lots of unique options; it went beyond the average spread at a hotel. I'm impressed.

--- Negative Example ---  
I went to Syracuse last winter to watch my team play in the Carrier Dome and chose this Econo Lodge because the prices were cheaper (because I found out later that it's in a pretty run-down part of town) and mainly cause we were at the better hotels were already booked. I drove up there and arrived pretty late, so I was very tired and ready to get some sleep before the game the next day. After pulling into the small parking lot, I realized after a driving around that all the Free Parking spaces were gone, so I had to drive down the street to a paid parking lot. Then, after carrying my luggage all the way to the hotel, I checked in and, in a sleepy haze, threw my stuff down in the room and went to sleep. Not more than two hours later, I woke up from incredibly itching on my arms and legs, so I quickly turned on the light and realized how awfully dirty my room was. Then I looked closer at my bed near the sheets and I literally could see many, if not hundreds, of tiny bedbugs crawling everywhere. I immediately ran to the desk, where no one was there of course, and after waiting for 30 minutes for someone to show up, I demanded another room. I then got the new room key, went up to it, and after opening the door, I saw that this new room was no more clean than the last. And yes, there were just as many bedbugs in this bed. I grabbed my stuff and walked back to my car and slept there for the rest of the night. Needless to say, not on I will I never visit this Econo Lodge again. I doubt I will ever even visit Syracuse again.

### Test Data (WAIT TILL DEADLINE)

This is the test dataset that you will need to use to report the results on. This set is the unseen dataset meaning, you are not in anyway suppose to look what is in this dataset. We will release this dataset on the last day of the assignment's deadline.

```
In [45]: ### RUN THIS ONLY ON DEADLINE ###
# Load the test data

from typing import List, Tuple, Any

def load_test_data(filepath: str) -> tuple[List[Any], list[Any]]:
    """Load the test data, producing a list of texts, labels

    Args:
        filepath (str): Path to the training file

    Returns:
        tuple([list[Any], list[Any]]): The texts and labels

    """
    lab_map = {'POS': 1, 'NEG': 0}
    texts = []
    labels = []
    with open(filepath, "r") as file:
        for line in file:
            idx, text, label = line.rstrip().split("\t")
            texts.append(text)
            lab_map[label].append(text)
        return texts, labels

test_texts, test_labels = load_test_data('./data/WN2-testset.txt')
```

### Task 1.1: Print the number of "positive" and "negative" samples (5 Points)

It is important to know the distribution of the training examples. More often than not, you will have to work with datasets that are not "balanced" with respect to the labels of the samples. For this task, print out the number of examples that have label = 1 and label = 0, respectively, in stdout or plot a pie chart.

```
In [41]: ### ENTER CODE HERE ###
import matplotlib.pyplot as plt
import numpy as np
# Note since we have them in 2 separate files,
# this can also be done with bash commands
def label_distribution(labels):
    """
    TODO: Replace the line 'raise NotImplementedError' with your code
    to print the labels distribution.

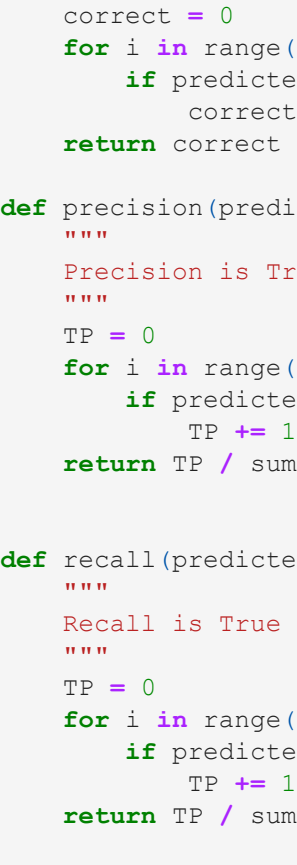
    """
    neg_count = pos_count = 0
    for num in labels:
        if num == 0:
            neg_count += 1
        elif num == 1:
            pos_count += 1
    print(f"The number of examples that have label = 0: ", neg_count)
    print(f"The number of examples that have label = 1: ", pos_count)

y = np.array([neg_count, pos_count])
mylabels = ["neg_count", "pos_count"]

plt.pie(y, labels = mylabels, autopct='%1.1f%%')
plt.show()

label_distribution(all_labels)
```

The number of examples that have label = 0: 94  
The number of examples that have label = 1: 95



### Task 1.2: Split Training and Development Sets (5 Points)

For the purpose of coming with the best parameters for the model you will have to split the dataset into training and development sets. Make sure the splits follow the same distribution.

```
In [47]: ### ENTER CODE HERE ###
import numpy as np
import pandas as pd
def split_dataset(texts, labels):
    """
    Split the dataset randomly into 80% training and 20% development set
    Make sure the splits have the same label distribution

    """
    indices = list(range(len(texts)))
    num_training_indices = int(0.8 * len(texts))
    np.random.shuffle(indices)
    train_indices = indices[num_training_indices:]
    dev_indices = indices[:num_training_indices]
    # Split the actual data
    train_texts, train_labels = texts.iloc[train_indices], labels.iloc[train_indices]
    dev_texts, dev_labels = texts.iloc[dev_indices], labels.iloc[dev_indices]
    return train_texts[0].values.tolist(), train_labels[0].values.tolist(), dev_texts[0].values.tolist(), dev_labels[0].values.tolist()

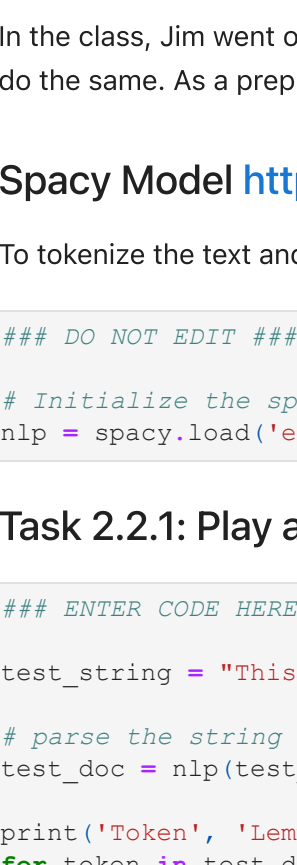
train_texts = pd.DataFrame(all_texts)
train_labels = pd.DataFrame(all_labels)

train_texts, train_labels, dev_texts, dev_labels = split_dataset(train_texts, train_labels)

print('Train Label Distribution:')
label_distribution(train_labels)

print('Dev Label Distribution:')
label_distribution(dev_labels)
```

Train Label Distribution:  
The number of examples that have label = 0: 82  
The number of examples that have label = 1: 69



### Task 1.3: Evaluation Metrics (10 Points)

Implement the evaluation metrics: Accuracy, Precision, Recall and F1 score

```
In [6]: ### ENTER CODE HERE ###
def accuracy(predicted_labels, true_labels):
    """
    Accuracy is correct predictions / all predictions

    """
    correct = 0
    for i in range(len(true_labels)):
        if predicted_labels[i] == true_labels[i]:
            correct += 1
    return correct / len(true_labels)

def precision(predicted_labels, true_labels):
    """
    Precision is True Positives / All Positives Predictions

    """
    TP = 0
    for i in range(len(true_labels)):
        if predicted_labels[i] == true_labels[i] and true_labels[i] == 1:
            TP += 1
    return TP / sum(predicted_labels)

def recall(predicted_labels, true_labels):
    """
    Recall is True Positives / All Positive Labels

    """
    TP = 0
    for i in range(len(true_labels)):
        if predicted_labels[i] == true_labels[i] and true_labels[i] == 1:
            TP += 1
    return TP / sum(true_labels)

def f1_score(predicted_labels, true_labels):
    """
    F1 score is the harmonic mean of precision and recall

    """
    prec = precision(predicted_labels, true_labels)
    reca = recall(predicted_labels, true_labels)
    return 2 * prec * reca / (prec + reca)
```

## Section 2: Baselines (Total: 20 Points)

It is important to come up with baselines for the classifications to compare the more complicated models with. The baselines are also useful as a debugging guide for your actual classification model. You will create two baselines:

1. Random Chance
2. Naive Bayes Classifier

### Task 2.1: Random Chance Classifier (5 Points)

A random chance classifier predicts the label according to the label's distribution. As an example, if the label 1 appears 70% of the times in the training set, you predict 70 out of 100 times the label 1 and label 0 30% of the times

```
In [7]: ### ENTER CODE HERE ###
def predict_random(train_labels, num_samples):
    """
    Predict random labels for num_samples samples.
    result = np.random.choice([0, 1], size=num_samples,
                             p=[1 - (sum(train_labels)/len(train_labels)), sum(train_labels)/len(train_labels)])

    """
    return result
```

### Task 2.2: Naive Bayes Classifier (Total: 10 Points)

In the class, Jim went over how to implement a Naive Bayes Classifier using the tokens in the training samples. In this task, you will do the same. As a preprocessing step, you might want to remove the stop words and lemmatize/stem the words of the texts.

#### Spacy Model <https://spacy.io>

To tokenize the text and help extract features from text, we will use the popular spaCy model

```
In [8]: ### DO NOT EDIT ###
# Initialize the spacy model
nlp = spacy.load('en_core_web_sm')
```

#### Task 2.2.1: Play around with spacy (0 Points)

```
In [9]: ### ENTER CODE HERE ###
test_string = "This is an amazing sentence"

# parse the string with spacy model
test_doc = nlp(test_string)

print('Tokens', 'Lemmas', 'Is_Stopped?')
for doc_features in test_doc:
    print(token, token.lemma_, token.is_stop)

Token Lemmas Is_Stopped?
This this True
is be True
an an True
amazing amazing False
sentence sentence False
```

#### Task 2.2.2: Preprocessing (5 Points)

Remove stopwords and lemmatize the words of a text

```
In [10]: ### ENTER CODE HERE ###
def pre_process(text: str) -> List[str]:
    """
    remove stopwords and lemmatize and return an array of lemmas

    """
    my_doc = nlp(text)

    lemma_nostop = []

    for token in my_doc:
        if token.is_stop == False:
            lemma_nostop.append(token.lemma_)
    return lemma_nostop

test_string = "This sentence needs to be lemmatized"

assert len({'sentence', 'need', 'lemmatize', 'lemmatiz'}.intersection(pre_process(test_string))) >= 3

print('All Test Cases Passed!')
```

All Test Cases Passed!

#### Task 2.2.3: The Naive Bayes Class (5 Points)

The standard way of implementing classifiers like Naive Bayes is to implement the two methods: "fit" and "predict". The fit method expects the training data along with labels, and the predict method predicts the labels for the provides texts of samples.

```
In [24]: ### ENTER CODE HERE ###
import operator
import math
class NaiveBayesClassifier:
    def __init__(self, num_classes):
        self.num_classes = num_classes
        self.label_words_count_pos = dict()
        self.label_words_count_neg = dict()

    def fit(self, texts, labels):
        """
        1. Group samples by their labels
        2. Preprocess each text
        3. Count the words of the text for each label

        """
        for i in range(len(texts)):
            vector = texts[i]
            class_value = int(labels[i])

            words = pre_process(vector)

            if class_value == 1:
                for word in words:
                    self.label_words_count_pos[word] = self.label_words_count_pos.get(word, 0) + 1
            elif class_value == 0:
                for word in words:
                    self.label_words_count_neg[word] = self.label_words_count_neg.get(word, 0) + 1

    def predict(self, texts):
        """
        1. Preprocess the texts
        2. Predict the class by using the likelihood with Bayes Method and Laplace Smoothing

        """
        pos_count_sum = sum(self.label_words_count_pos.values())
        neg_count_sum = sum(self.label_words_count_neg.values())

        label_words_prob_pos = self.label_words_count_pos.copy()
        label_words_prob_neg = self.label_words_count_neg.copy()

        label_words_prob_pos = {key: value / pos_count_sum for key, value in label_words_prob_pos.items()}
        label_words_prob_neg = {key: value / neg_count_sum for key, value in label_words_prob_neg.items()}

        result = []

        for i in range(len(texts)):
            vector = texts[i]
            words = pre_process(vector)

            pos_score = 0
            neg_score = 0

            for word in words:
                if word in self.label_words_count_pos:
                    pos_score += math.log(label_words_prob_pos[word])
                elif word not in self.label_words_count_pos:
                    pos_score = math.log(float(1/pos_count_sum))

                if word in self.label_words_count_neg:
                    neg_score += math.log(label_words_prob_neg[word])
                elif word not in self.label_words_count_neg:
                    neg_score = math.log(float(1/neg_count_sum))

            if pos_score > neg_score:
                result.append(1)
            else:
                result.append(0)

        return result
```

#### Task 2.3: Baseline Results (5 Points)

Since there is not hyperparameter-tuning required for the baselines, we can use the entirety of the training set (no need to split the dataset into train and development). Report the results you achieve with the two baselines by running the following cell:

```
In [12]: ### DO NOT EDIT ###

### DEV SET RESULTS
def predict_random = predict_random(train_labels, num_samples=len(dev_labels))
print('Random Chance F1:', f1_score(testset_prediction_random, dev_labels))

naive_bayes_classifier = NaiveBayesClassifier(num_classes=2)
naive_bayes_classifier.fit(train_texts, train_labels)
testset_predictions_nb = naive_bayes_classifier.predict(dev_texts)
print('Naive Bayes F1:', f1_score(testset_predictions_nb, dev_labels))

Random Chance F1: 0.5945945945945946
Naive Bayes F1: 0.9090909090909091
```

```
In [27]: ### DO NOT EDIT ###
### RUN THIS ONLY ON DEADLINE ###
### TEST SET RESULTS
testset_prediction_random = predict_random(all_labels, num_samples=len(test_labels))
print('Random Chance F1:', f1_score(testset_prediction_random, test_labels))

naive_bayes_classifier = NaiveBayesClassifier(num_classes=2)
naive_bayes_classifier.fit(all_texts, all_labels)
testset_predictions_nb = naive_bayes_classifier.predict(test_texts)
print('Naive Bayes F1:', f1_score(testset_predictions_nb, test_labels))

Random Chance F1: 0.4912280707543866
Naive Bayes F1: 0.9019607843137256
```

## Section 3: Logistic Regression on Features (Total: 60 Points)

Now let's try building a logistic regression based classifier on hand-engineered features.

The following tasks are going to be the implementation of the components required in building a Logistic Regressor.

### Task 3.0: Feature Extraction (20 points)

This is perhaps the most challenging part of this assignment. In the class, we went over how to featurize text for a classification system for sentiment analysis. In this assignment, you should implement and build upon this to accurately classify the hotel reviews.

This task requires a thorough understanding of the dataset to answer the important question, "What is in the data?". Please go through some of the datapoints and convert the signals that you think might help in identifying "sentiment" as features.

Please refer to the section in Jim's book that illustrates the process of feature engineering for this task. We have attached an image of the table below:

Var	Definition	Value in Fig. 5.2
$x_1$	$\text{count}(\text{positive lexicon}) \in \text{doc}$	3
$x_2$	$\text{count}(\text{negative lexicon}) \in \text{doc}$	2
$x_3$	$\begin{cases} 1 & \text{if "no"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	1
$x_4$	$\text{count}(\text{1st and 2nd pronouns}) \in \text{doc}$	3
$x_5$	$\begin{cases} 1 & \text{if "I"} \in \text{doc} \\ 0 & \text{otherwise} \end{cases}$	0
$x_6$	$\log(\text{word count of doc})$	$\ln(64) = 4.15$

Please use the files with positive and negative words attached in the assignment: `positive_words.txt` and `negative_words.txt`

```
In [13]: from util import load_train_data
pos_datapath = "data/positive_words.txt"
neg_datapath = "data/negative_words.txt"

my_file = open(pos_datapath, "r")
content = my_file.read()
pos_lex = content.split("\n")
my_file.close()

my_file = open(neg_datapath, "r")
content = my_file.read()
neg_lex = content.split("\n")
my_file.close()

def make_test_feature_1(text: spacy.tokens.doc.Doc):
    count = 0
    for t in text:
        tem = t.lemma_
        if t.lemma_ in pos_lex:
            count += 1
    return count

def make_test_feature_2(text: spacy.tokens.doc.Doc):
    count = 0
    for t in text:
        if t.lemma_ in neg_lex:
            count += 1
    return count

def make_test_feature_3(text: spacy.tokens.doc.Doc):
    count = 0
    for t in text:
        if t.lemma_ == 'no':
            count += 1
    return count

def make_test_feature_4(text: spacy.tokens.doc.Doc):
    count = 0
    pronouns = ['I', 'me', 'my', 'My', 'mine', 'myself', 'You', 'you', 'your', 'Your', 'yours', 'yourself']
    for t in text:
        if t.lemma_ in pronouns:
            count += 1
    return count

def make_test_feature_5(text: spacy.tokens.doc.Doc):
    count = 0
    for t in text:
        if t.lemma_ == 'I':
            count += 1
    return count

def make_test_feature_6(text: spacy.tokens.doc.Doc):
    count = 0
    for t in text:
        count += 1
    return np.log(count)

def extract_features(text: spacy.tokens.doc.Doc):
    features = []
    # TODO: Replace this with your own feature extraction functions.
    # TODO: add more features to the feature vector
    features.append(make_test_feature_1(text))
    features.append(make_test_feature_2(text))
    features.append(make_test_feature_3(text))
    features.append(make_test_feature_4(text))
    features.append(make_test_feature_5(text))
    features.append(make_test_feature_6(text))

    return features
```

```
In [14]: ### ENTER CODE HERE ###
### DO NOT CHANGE THE SIGNATURE OF THE function THOUGH ###
def featurize_data(texts, labels):
    """
    extract_features(doc) for doc in nlp.pipe(texts)

    """
    return torch.FloatTensor(features), torch.FloatTensor(labels)
```

### Task 3.0.2: Feature Scaling (10 Points)

In this task we will use the data normalization technique to ensure the scales of the feature are consistent. After featurizing the dataset, we need to call the following function before passing it to the classifier

#### Normalization Formula

$$X' = \frac{X - X_{min}}{X_{max} - X_{min}}$$

```
In [15]: ### ENTER CODE HERE ###
def normalize(features: torch.Tensor) -> torch.Tensor:
    """
    return the features transformed by the above formula of normalization

    """
    feature_min = torch.min(features)
    feature_max = torch.max(features)
    for row in range(features.shape[0]):
        for col in range(features.shape[1]):
            features[row][col] = (features[row][col].item() - feature_min) / (feature_max - feature_min)

    return features
```

## Training a Logistic Regression Classifier (Total: 30 Points)

In this section, you will implement the components needed to train the binary classifier using logistic regression

### Here we define our pytorch logistic regression classifier (DO NOT EDIT THIS)

```
In [16]: class SentimentClassifier(torch.nn.Module):
    def __init__(self, input_dim: int):
        super().__init__()
        # We force output to be one, since we are doing binary logistic regression
        self.output_size = 1
        self.coefficients = torch.nn.Linear(input_dim, self.output_size)
        # Initialize weights. Note that this is not strictly necessary,
        # but you should test different initializations per lecture
        initialize_weights(self.coefficients)

    def forward(self, features: torch.Tensor):
        # We predict a number by multiplying by the coefficients
        # and then take the sigmoid to turn the score as logits
        return torch.sigmoid(self.coefficients(features))
```

#### Task 3.1: Initialize the weights. (5 Points)

Initialization of the parameters is an important step to ensure the SGD algorithm converges to a global optimum. Typically, we need to try different initialization methods and compare the accuracy we achieve for the development set. In this task, implement the function that initializes the parameters to ...

```
In [17]: ### ENTER CODE HERE ###
import math
def initialize_weights(coefficients):
    """
    TODO: Replace the line 'raise NotImplementedError' with your code.
    Initialize the weights of the coefficients by assigning the parameter
    coefficients.weights.data = ...

    """
    return torch.nn.init.xavier_uniform_(coefficients.weight)
```

Let's build a training function similar to the linear regressor from the tutorial

#### Task 3.2: Logistic Loss Function (10 Points)

```
In [18]: ### ENTER CODE HERE ###
def logistic_loss(prediction: torch.Tensor, label: torch.Tensor) -> torch.Tensor:
    """
    TODO: Implement the logistic loss function between a prediction and label.

    """
    ll = torch.nn.BCEWithLogitsLoss()
    return ll(prediction, label)
```

#### Task 3.3: Create an SGD optimizer (0 Points)

We have already provided the implementation of how to create the SGD optimizer

You may try different optimizers referring to the docs provided

```
In [19]: ### ENTER CODE HERE ###
def make_optimizer(model, learning_rate) -> torch.optim:
    """
    Returns an Stochastic Gradient Descent Optimizer
    See here for algorithms you can import: https://pytorch.org/docs/stable/optim.html
    """
    return torch.optim.SGD(model.parameters(), learning_rate)
```

#### Task 3.5: Converting Logits into Predictions (5 Points)

```
In [20]: ### ENTER CODE HERE ###
def predict(model, features):
    """
    TODO: Replace the line 'raise NotImplementedError'
    with the logic of converting the logits into prediction labels (0, 1)

    """
    logits = model(features)
    return torch.round(logits)
```

### Training Function (DO NOT EDIT THIS)

```
In [21]: ### DO NOT EDIT ###
from tqdm.autonotebook import tqdm
import random

def training_loop(
    num_epochs,
    batch_size,
    train_features,
    train_labels,
    dev_features,
    dev_labels,
    model,
    optimizer,
):
    samples = list(zip(train_features, train_labels))
    random.shuffle(samples)
    batches = []
    for i in range(0, len(samples), batch_size):
        batches.append(samples[i:i+batch_size])
    print("Training...")
    for i in range(num_epochs):
        losses = []
        for batch in tqdm(batches):
            # Empty the dynamic computation graph
            features, labels = zip(*batch)
            features = torch.stack(features)
            labels = torch.stack(labels)
            optimizer.zero_grad()
            # Run the model
            logits = model(features)
            # compute loss
            loss = logistic_loss(torch.squeeze(logits), labels)
            # In this logistic regression example,
            # this entails computing a single gradient
            loss.backward()
            # Backpropagate the loss through our model

            # Update our coefficients in the direction of the gradient.
            optimizer.step()
            # For logging
            losses.append(loss.item())

        # Estimate the F1 score for the development set
        dev_f1 = f1_score(predict(model, dev_features), dev_labels)
        print(f"epoch {i}, loss: {sum(losses)/len(losses)}")
        print(f"Dev F1 {dev_f1}")

    # Return the trained model
    return model
```

#### Task 3.6: Train the classifier (10 Points)

Run the following cell to train a logistic regressor on your hand-engineered features.

```
In [52]: num_epochs = 500

train_features, train_labels_tensor = featurize_data(train_texts, train_labels)
train_features = normalize(train_features)
dev_features, dev_labels_tensor = featurize_data(dev_texts, dev_labels)
dev_features = normalize(dev_features)
model = SentimentClassifier(train_features.shape[1])
optimizer = make_optimizer(model, learning_rate=0.5)

trained_model = training_loop(
    num_epochs,
    train_features,
    train_labels_tensor,
    dev_features,
    dev_labels_tensor,
    model,
    optimizer,
)
```



Training...  
epoch 0, Loss: 0.7417197108268730  
Dev F1 tensor(0.3226)  
epoch 1, Loss: 0.733013854446411  
Dev F1 tensor(0.3226)  
epoch 2, Loss: 0.7263218641281128  
Dev F1 tensor(0.3226)  
epoch 3, Loss: 0.7260681300393223  
Dev F1 tensor(0.3226)  
epoch 4, Loss: 0.7161434412002563  
Dev F1 tensor(0.3226)  
epoch 5, Loss: 0.7125257790886504  
Dev F1 tensor(0.3226)  
epoch 6, Loss: 0.7096132866437988  
Dev F1 tensor(0.3226)  
epoch 7, Loss: 0.7072575688362122  
Dev F1 tensor(0.3226)  
epoch 8, Loss: 0.705322003464563  
Dev F1 tensor(0.3226)  
epoch 9, Loss: 0.7037158846855164  
Dev F1 tensor(0.3226)  
epoch 10, Loss: 0.702377825927734  
Dev F1 tensor(0.3226)  
epoch 11, Loss: 0.70122362352594  
Dev F1 tensor(0.3226)  
epoch 12, Loss: 0.7002422451979262  
Dev F1 tensor(0.3226)  
epoch 13, Loss: 0.6993921220320581  
Dev F1 tensor(0.3226)  
epoch 14, Loss: 0.6986488223075866  
Dev F1 tensor(0.3226)  
epoch 15, Loss: 0.697993278503418  
Dev F1 tensor(0.3226)  
epoch 16, Loss: 0.6974104404449463  
Dev F1 tensor(0.3226)  
epoch 17, Loss: 0.6968688267993927  
Dev F1 tensor(0.3226)  
epoch 18, Loss: 0.696417099237442  
Dev F1 tensor(0.3226)  
epoch 19, Loss: 0.6959891319274902  
Dev F1 tensor(0.3226)  
epoch 20, Loss: 0.695597982999649  
Dev F1 tensor(0.3226)  
epoch 21, Loss: 0.6952381193637848  
Dev F1 tensor(0.3226)  
epoch 22, Loss: 0.6949045412319713  
Dev F1 tensor(0.3226)  
epoch 23, Loss: 0.694596083130243  
Dev F1 tensor(0.3226)  
epoch 24, Loss: 0.6943046981563581  
Dev F1 tensor(0.3226)  
epoch 25, Loss: 0.6940254943275452  
Dev F1 tensor(0.3226)  
epoch 26, Loss: 0.6937793254852295  
Dev F1 tensor(0.3226)  
epoch 27, Loss: 0.6935365200042725  
Dev F1 tensor(0.3226)  
epoch 28, Loss: 0.693295456638362  
Dev F1 tensor(0.3226)  
epoch 29, Loss: 0.6930846154689789  
Dev F1 tensor(0.3226)  
epoch 30, Loss: 0.6928727149963378  
Dev F1 tensor(0.3226)  
epoch 31, Loss: 0.6926686227321625  
Dev F1 tensor(0.3226)  
epoch 32, Loss: 0.692471349239494  
Dev F1 tensor(0.3226)  
epoch 33, Loss: 0.692280660133361  
Dev F1 tensor(0.3226)  
epoch 34, Loss: 0.692093851425934  
Dev F1 tensor(0.3226)  
epoch 35, Loss: 0.6919121026992798  
Dev F1 tensor(0.3226)  
epoch 36, Loss: 0.691734117269516  
Dev F1 tensor(0.3226)  
epoch 37, Loss: 0.6915593326091767  
Dev F1 tensor(0.3226)  
epoch 38, Loss: 0.6913872182639232  
Dev F1 tensor(0.3226)  
epoch 39, Loss: 0.6912172257900238  
Dev F1 tensor(0.3226)  
epoch 40, Loss: 0.6910490959615387  
Dev F1 tensor(0.3226)  
epoch 41, Loss: 0.690882086738452  
Dev F1 tensor(0.3226)  
epoch 42, Loss: 0.6907160639762878  
Dev F1 tensor(0.3226)  
epoch 43, Loss: 0.6905506193637848  
Dev F1 tensor(0.3226)  
epoch 44, Loss: 0.690385392886474  
Dev F1 tensor(0.3226)  
epoch 45, Loss: 0.690220996875762  
Dev F1 tensor(0.3226)  
epoch 46, Loss: 0.6900543987751007  
Dev F1 tensor(0.3226)  
epoch 47, Loss: 0.6898880362016081  
Dev F1 tensor(0.3226)  
epoch 48, Loss: 0.6897207081317902  
Dev F1 tensor(0.3226)  
epoch 49, Loss: 0.6895521819591522  
Dev F1 tensor(0.3226)  
epoch 50, Loss: 0.689382431564331  
Dev F1 tensor(0.3226)  
epoch 51, Loss: 0.68921059660734  
Dev F1 tensor(0.3226)  
epoch 52, Loss: 0.6890370845794678  
Dev F1 tensor(0.3226)  
epoch 53, Loss: 0.688861483335495  
Dev F1 tensor(0.3226)  
epoch 54, Loss: 0.688683545899447  
Dev F1 tensor(0.3226)  
epoch 55, Loss: 0.6885031044483185  
Dev F1 tensor(0.3226)  
epoch 56, Loss: 0.6883199572563171  
Dev F1 tensor(0.3226)  
epoch 57, Loss: 0.6881339848041534  
Dev F1 tensor(0.3226)  
epoch 58, Loss: 0.68794949327127838  
Dev F1 tensor(0.3226)  
epoch 59, Loss: 0.687752777338028  
Dev F1 tensor(0.3226)  
epoch 60, Loss: 0.6875571846861975  
Dev F1 tensor(0.3226)  
epoch 61, Loss: 0.6873581647872925  
Dev F1 tensor(0.3226)  
epoch 62, Loss: 0.6871555209159851  
Dev F1 tensor(0.3226)  
epoch 63, Loss: 0.686949636257727  
Dev F1 tensor(0.3226)  
epoch 64, Loss: 0.686739857769013  
Dev F1 tensor(0.3226)  
epoch 65, Loss: 0.6865248620510102  
Dev F1 tensor(0.3226)  
epoch 66, Loss: 0.6863066971302032  
Dev F1 tensor(0.3226)  
epoch 67, Loss: 0.6860845148563385  
Dev F1 tensor(0.3226)  
epoch 68, Loss: 0.685861602573395  
Dev F1 tensor(0.3226)  
epoch 69, Loss: 0.6856276154518127  
Dev F1 tensor(0.3226)  
epoch 70, Loss: 0.685392896306872  
Dev F1 tensor(0.3226)  
epoch 71, Loss: 0.685159373397827  
Dev F1 tensor(0.3226)  
epoch 72, Loss: 0.6849107801914215  
Dev F1 tensor(0.3226)  
epoch 73, Loss: 0.6846634387969971  
Dev F1 tensor(0.3226)  
epoch 74, Loss: 0.6844119250773483  
Dev F1 tensor(0.3226)  
epoch 75, Loss: 0.6841563165187836  
Dev F1 tensor(0.3226)  
epoch 76, Loss: 0.683896727256774  
Dev F1 tensor(0.3226)  
epoch 77, Loss: 0.683630473423005  
Dev F1 tensor(0.3226)  
epoch 78, Loss: 0.6833577459353  
Dev F1 tensor(0.3226)  
epoch 79, Loss: 0.683094364044783  
Dev F1 tensor(0.3226)  
epoch 80, Loss: 0.6828195035457612  
Dev F1 tensor(0.3226)  
epoch 81, Loss: 0.6825411736965179  
Dev F1 tensor(0.3226)  
epoch 82, Loss: 0.682259476184845  
Dev F1 tensor(0.3226)  
epoch 83, Loss: 0.6819745838642121  
Dev F1 tensor(0.3226)  
epoch 84, Loss: 0.681686693904114  
Dev F1 tensor(0.3226)  
epoch 85, Loss: 0.6813959598541259  
Dev F1 tensor(0.3226)  
epoch 86, Loss: 0.6811024785041809  
Dev F1 tensor(0.3226)  
epoch 87, Loss: 0.680805593242445  
Dev F1 tensor(0.3226)  
epoch 88, Loss: 0.6805082798004151  
Dev F1 tensor(0.3226)  
epoch 89, Loss: 0.6802078723907471  
Dev F1 tensor(0.3226)  
epoch 90, Loss: 0.6799054920673371  
Dev F1 tensor(0.3226)  
epoch 91, Loss: 0.679601329565483  
Dev F1 tensor(0.3226)  
epoch 92, Loss: 0.6792955935001374  
Dev F1 tensor(0.3226)  
epoch 93, Loss: 0.6789883971214294  
Dev F1 tensor(0.3226)  
epoch 94, Loss: 0.678679925203233  
Dev F1 tensor(0.3226)  
epoch 95, Loss: 0.6783703446388245  
Dev F1 tensor(0.3226)  
epoch 96, Loss: 0.6780597984790802  
Dev F1 tensor(0.3226)  
epoch 97, Loss: 0.6777484655380249  
Dev F1 tensor(0.3226)  
epoch 98, Loss: 0.677432492624613  
Dev F1 tensor(0.3226)  
epoch 99, Loss: 0.6771238088607788  
Dev F1 tensor(0.3226)  
epoch 100, Loss: 0.67681081283501342  
Dev F1 tensor(0.3226)  
epoch 101, Loss: 0.6764975190162659  
Dev F1 tensor(0.3226)  
epoch 102, Loss: 0.6761840164661408  
Dev F1 tensor(0.3226)  
epoch 103, Loss: 0.675870328785791  
Dev F1 tensor(0.3226)  
epoch 104, Loss: 0.6755567848682403  
Dev F1 tensor(0.3226)  
epoch 105, Loss: 0.6752432346343994  
Dev F1 tensor(0.3226)  
epoch 106, Loss: 0.6749298751354218  
Dev F1 tensor(0.3226)  
epoch 107, Loss: 0.6746167540505232  
Dev F1 tensor(0.3226)  
epoch 108, Loss: 0.674303936395813  
Dev F1 tensor(0.3226)  
epoch 109, Loss: 0.6739915132522583  
Dev F1 tensor(0.3226)  
epoch 110, Loss: 0.6736795365810394  
Dev F1 tensor(0.3226)  
epoch 111, Loss: 0.6733680650888367  
Dev F1 tensor(0.3226)  
epoch 112, Loss: 0.6730570979196579  
Dev F1 tensor(0.3226)  
epoch 113, Loss: 0.6727467536926269  
Dev F1 tensor(0.3226)  
epoch 114, Loss: 0.672437059879303  
Dev F1 tensor(0.3226)  
epoch 115, Loss: 0.6721280515193939  
Dev F1 tensor(0.3226)  
epoch 116, Loss: 0.6718197584152221  
Dev F1 tensor(0.3226)  
epoch 117, Loss: 0.6715122640132904  
Dev F1 tensor(0.3226)  
epoch 118, Loss: 0.671205444717407  
Dev F1 tensor(0.3226)  
epoch 119, Loss: 0.67089965347534  
Dev F1 tensor(0.3226)  
epoch 120, Loss: 0.670594662479022  
Dev F1 tensor(0.3226)  
epoch 121, Loss: 0.670290541648647  
Dev F1 tensor(0.3226)  
epoch 122, Loss: 0.6699873387813569  
Dev F1 tensor(0.3226)  
epoch 123, Loss: 0.6696851372718811  
Dev F1 tensor(0.3226)  
epoch 124, Loss: 0.66938329832077  
Dev F1 tensor(0.3226)  
epoch 125, Loss: 0.669083547592163  
Dev F1 tensor(0.3226)  
epoch 126, Loss: 0.6687842726707458  
Dev F1 tensor(0.3226)  
epoch 127, Loss: 0.6684866052715412  
Dev F1 tensor(0.3226)  
epoch 128, Loss: 0.66818882299042  
Dev F1 tensor(0.3226)  
epoch 129, Loss: 0.6678927361965179  
Dev F1 tensor(0.3226)  
epoch 130, Loss: 0.6675977051258087  
Dev F1 tensor(0.3226)  
epoch 131, Loss: 0.6673037707805634  
Dev F1 tensor(0.3226)  
epoch 132, Loss: 0.6670109570026398  
Dev F1 tensor(0.3226)  
epoch 133, Loss: 0.66671923995018  
Dev F1 tensor(0.3226)  
epoch 134, Loss: 0.666428643650422  
Dev F1 tensor(0.3226)  
epoch 135, Loss: 0.6661392629146576  
Dev F1 tensor(0.3226)  
epoch 136, Loss: 0.665850967168808  
Dev F1 tensor(0.3226)  
epoch 137, Loss: 0.665563875437829  
Dev F1 tensor(0.3226)  
epoch 138, Loss: 0.6652779519557953  
Dev F1 tensor(0.3226)  
epoch 139, Loss: 0.6649932205677033  
Dev F1 tensor(0.3226)  
epoch 140, Loss: 0.6647096574306488  
Dev F1 tensor(0.3226)  
epoch 141, Loss: 0.6644272863864898  
Dev F1 tensor(0.3226)  
epoch 142, Loss: 0.6641461193561554  
Dev F1 tensor(0.3226)  
epoch 143, Loss: 0.6638661861419678  
Dev F1 tensor(0.3226)  
epoch 144, Loss: 0.6635874330997467  
Dev F1 tensor(0.3226)  
epoch 145, Loss: 0.6633099019527435  
Dev F1 tensor(0.3226)  
epoch 146, Loss: 0.6630335986414228  
Dev F1 tensor(0.3226)  
epoch 147, Loss: 0.6627584755420685  
Dev F1 tensor(0.3226)  
epoch 148, Loss: 0.66248459815979  
Dev F1 tensor(0.3226)  
epoch 149, Loss: 0.662211936712265  
Dev F1 tensor(0.3226)  
epoch 150, Loss: 0.6619405150413513  
Dev F1 tensor(0.3226)  
epoch 151, Loss: 0.6616703152656556  
Dev F1 tensor(0.3226)  
epoch 152, Loss: 0.661401343345642  
Dev F1 tensor(0.3226)  
epoch 153, Loss: 0.66113393208465  
Dev F1 tensor(0.3226)  
epoch 154, Loss: 0.6608670771213978  
Dev F1 tensor(0.3226)  
epoch 155, Loss: 0.6606017649173737  
Dev F1 tensor(0.3226)  
epoch 156, Loss: 0.6603376924991607  
Dev F1 tensor(0.3226)  
epoch 157, Loss: 0.6600748479366303  
Dev F1 tensor(0.3226)  
epoch 158, Loss: 0.6598132073879241  
Dev F1 tensor(0.3226)  
epoch 159, Loss: 0.6595528185367584  
Dev F1 tensor(0.3226)  
epoch 160, Loss: 0.6592936396598816  
Dev F1 tensor(0.3226)  
epoch 161, Loss: 0.6590356707529397  
Dev F1 tensor(0.3226)  
epoch 162, Loss: 0.6587789237499238  
Dev F1 tensor(0.3226)  
epoch 163, Loss: 0.6585233867168426  
Dev F1 tensor(0.3226)  
epoch 164, Loss: 0.658269077539444  
Dev F1 tensor(0.3226)  
epoch 165, Loss: 0.6580159549494764  
Dev F1 tensor(0.3226)  
epoch 166, Loss: 0.65776459350191  
Dev F1 tensor(0.3226)  
epoch 167, Loss: 0.6575133621692657  
Dev F1 tensor(0.3226)  
epoch 168, Loss: 0.6572638154029846  
Dev F1 tensor(0.3226)  
epoch 169, Loss: 0.6570155203342438  
Dev F1 tensor(0.3226)  
epoch 170, Loss: 0.656768411397934  
Dev F1 tensor(0.3226)  
epoch 171, Loss: 0.6565224707124618  
Dev F1 tensor(0.3226)  
epoch 172, Loss: 0.6562776863574982  
Dev F1 tensor(0.3226)  
epoch 173, Loss: 0.656034117937088  
Dev F1 tensor(0.3226)  
epoch 174, Loss: 0.655791729686444  
Dev F1 tensor(0.3226)  
epoch 175, Loss: 0.6555494918098449  
Dev F1 tensor(0.3226)  
epoch 176, Loss: 0.6553104400634766  
Dev F1 tensor(0.3226)  
epoch 177, Loss: 0.6550715446472168  
Dev F1 tensor(0.3226)  
epoch 178, Loss: 0.6548337996060102  
Dev F1 tensor(0.3226)  
epoch 179, Loss: 0.6545972168445587  
Dev F1 tensor(0.3226)  
epoch 180, Loss: 0.6543617725372315  
Dev F1 tensor(0.3226)  
epoch 181, Loss: 0.6541274666786194  
Dev F1 tensor(0.3226)  
epoch 182, Loss: 0.6538943231105805  
Dev F1 tensor(0.3226)  
epoch 183, Loss: 0.6536622885889343  
Dev F1 tensor(0.3226)  
epoch 184, Loss: 0.6534313976746679  
Dev F1 tensor(0.3226)  
epoch 185, Loss: 0.6532016277313233  
Dev F1 tensor(0.3226)  
epoch 186, Loss: 0.652972960472107  
Dev F1 tensor(0.3226)  
epoch 187, Loss: 0.652745449549993  
Dev F1 tensor(0.3226)  
epoch 188, Loss: 0.6525190055370331  
Dev F1 tensor(0.3226)  
epoch 189, Loss: 0.6522936701774598  
Dev F1 tensor(0.3226)  
epoch 190, Loss: 0.652069434642792  
Dev F1 tensor(0.3226)  
epoch 191, Loss: 0.6518463015556335  
Dev F1 tensor(0.3226)  
epoch 192, Loss: 0.651624262329163  
Dev F1 tensor(0.3226)  
epoch 193, Loss: 0.6514032781124115  
Dev F1 tensor(0.3226)  
epoch 194, Loss: 0.6511834144592286  
Dev F1 tensor(0.3226)  
epoch 195, Loss: 0.6509645760059357  
Dev F1 tensor(0.3226)  
epoch 196, Loss: 0.6507468163967133  
Dev F1 tensor(0.3226)  
epoch 197, Loss: 0.6505301296710968  
Dev F1 tensor(0.3226)  
epoch 198, Loss: 0.6503148602767639  
Dev F1 tensor(0.3226)  
epoch 199, Loss: 0.650099152660369  
Dev F1 tensor(0.3226)  
epoch 200, Loss: 0.6498863518238067  
Dev F1 tensor(0.3226)  
epoch 201, Loss: 0.6496738612651825  
Dev F1 tensor(0.3226)  
epoch 202, Loss: 0.6494623959064484  
Dev F1 tensor(0.3226)  
epoch 203, Loss: 0.6492519378662109  
Dev F1 tensor(0.3226)  
epoch 204, Loss: 0.6490424990653991  
Dev F1 tensor(0.3226)  
epoch 205, Loss: 0.648834091424942  
Dev F1 tensor(0.3226)  
epoch 206, Loss: 0.6486267931820526  
Dev F1 tensor(0.3226)  
epoch 207, Loss: 0.6484202980995178  
Dev F1 tensor(0.3226)  
epoch 208, Loss: 0.6482148885726928  
Dev F1 tensor(0.3226)  
epoch 209, Loss: 0.6480104982852936  
Dev F1 tensor(0.3226)  
epoch 210, Loss: 0.6478070795536042  
Dev F1 tensor(0.3226)  
epoch 211, Loss: 0.6476046562194824  
Dev F1 tensor(0.3226)  
epoch 212, Loss: 0.6474031925201416  
Dev F1 tensor(0.3226)  
epoch 213, Loss: 0.6472027363292975  
Dev F1 tensor(0.3226)  
epoch 214, Loss: 0.6470032155513763  
Dev F1 tensor(0.3226)  
epoch 215, Loss: 0.6468046605587006  
Dev F1 tensor(0.3226)  
epoch 216, Loss: 0.6466070652008057  
Dev F1 tensor(0.3226)  
epoch 217, Loss: 0.6464104056358337  
Dev F1 tensor(0.3226)  
epoch 218, Loss: 0.6462147116661072  
Dev F1 tensor(0.3226)  
epoch 219, Loss: 0.646019356079102  
Dev F1 tensor(0.3226)  
epoch 220, Loss: 0.6458261072636565  
Dev F1 tensor(0.3226)  
epoch 221, Loss: 0.6456332087516785  
Dev F1 tensor(0.3226)  
epoch 222, Loss: 0.6454401228315214  
Dev F1 tensor(0.3226)  
epoch 223, Loss: 0.6452501595020295  
Dev F1 tensor(0.3226)  
epoch 224, Loss: 0.6450600266456604  
Dev F1 tensor(0.3226)  
epoch 225, Loss: 0.6448707699757695  
Dev F1 tensor(0.3226)  
epoch 226, Loss: 0.6446824312210083  
Dev F1 tensor(0.3226)  
epoch 227, Loss: 0.6444949865341186  
Dev F1 tensor(0.3226)  
epoch 228, Loss: 0.6443084418773651  
Dev F1 tensor(0.3226)  
epoch 229, Loss: 0.6441227555274963  
Dev F1 tensor(0.3226)  
epoch 230, Loss: 0.6439379751682281  
Dev F1 tensor(0.3226)  
epoch 231, Loss: 0.6437540533158448  
Dev F1 tensor(0.3226)  
epoch 232, Loss: 0.6435710193726685  
Dev F1 tensor(0.3226)  
epoch 233, Loss: 0.643388837591525  
Dev F1 tensor(0.3226)  
epoch 234, Loss: 0.6432075083259768  
Dev F1 tensor(0.3226)  
epoch 235, Loss: 0.6430270314216614  
Dev F1 tensor(0.3226)  
epoch 236, Loss: 0.6428473946423373  
Dev F1 tensor(0.3226)  
epoch 237, Loss: 0.642668406135559  
Dev F1 tensor(0.3226)  
epoch 238, Loss: 0.642490702867508  
Dev F1 tensor(0.3226)  
epoch 239, Loss: 0.642313876655579  
Dev F1 tensor(0.3226)  
epoch 240, Loss: 0.6421373008681701  
Dev F1 tensor(0.3226)  
epoch 241, Loss: 0.6419618725776672  
Dev F1 tensor(0.3226)  
epoch 242, Loss: 0.641787248848688  
Dev F1 tensor(0.3226)  
epoch 243, Loss: 0.6416134059429168  
Dev F1 tensor(0.3226)  
epoch 244, Loss: 0.641440395909918  
Dev F1 tensor(0.3226)  
epoch 245, Loss: 0.6412687996882357  
Dev F1 tensor(0.3226)  
epoch 246, Loss: 0.6410968005657196  
Dev F1 tensor(0.3226)  
epoch 247, Loss: 0.6409261763095856  
Dev F1 tensor(0.3226)  
epoch 248, Loss: 0.64075634795227  
Dev F1 tensor(0.3226)  
epoch 249, Loss: 0.6405873006621796  
Dev F1 tensor(0.3226)  
epoch 250, Loss: 0.640419480709057  
Dev F1 tensor(0.3226)  
epoch 251, Loss: 0.6402515649795533  
Dev F1 tensor(0.3226)  
epoch 252, Loss: 0.640084848274519  
Dev F1 tensor(0.3226)  
epoch 253, Loss: 0.639918914570618  
Dev F1 tensor(0.3226)  
epoch 254, Loss: 0.639753713840668  
Dev F1 tensor(0.3226)  
epoch 255, Loss: 0.6395892386168091  
Dev F1 tensor(0.3226)  
epoch 256, Loss: 0.6394254055355072  
Dev F1 tensor(0.3226)  
epoch 257, Loss: 0.639262404742268  
Dev F1 tensor(0.3226)  
epoch 258, Loss: 0.63910456237793  
Dev F1 tensor(0.3226)  
epoch 259, Loss: 0.6389490551364898  
Dev F1 tensor(0.3226)  
epoch 260, Loss: 0.6387978252513885  
Dev F1 tensor(0.3226)  
epoch 261, Loss: 0.6386482904243469  
Dev F1 tensor(0.3226)  
epoch 262, Loss: 0.638499029714491  
Dev F1 tensor(0.3226)  
epoch 263, Loss: 0.6383504729797401  
Dev F1 tensor(0.3226)  
epoch 264, Loss: 0.6382034976753234  
Dev F1 tensor(0.3226)  
epoch 265, Loss: 0.6380595497713089  
Dev F1 tensor(0.3226)  
epoch 266, Loss: 0.6379179272438049  
Dev F1 tensor(0.3226)  
epoch 267, Loss: 0.6377793408325806  
Dev F1 tensor(0.3226)  
epoch 268, Loss: 0.6376433403491974  
Dev F1 tensor(0.3226)  
epoch 269, Loss: 0.6375093623928739  
Dev F1 tensor(0.3226)  
epoch 270, Loss: 0.6373780750705719  
Dev F1 tensor(0.3226)  
epoch 271, Loss: 0.6372493593544006  
Dev F1 tensor(0.3226)  
epoch 272, Loss: 0.63712375447464  
Dev F1 tensor(0.3226)  
epoch 273, Loss: 0.636999769393423  
Dev F1 tensor(0.3226)  
epoch 274, Loss: 0.6368782639938355  
Dev F1 tensor(0.3226)  
epoch 275, Loss: 0.636759438466791  
Dev F1 tensor(0.3226)  
epoch 276, Loss: 0.6366430810907368  
Dev F1 tensor(0.3226)  
epoch 277, Loss: 0.63652813882749281  
Dev F1 tensor(0.3226)  
epoch 278, Loss: 0.6364153964944428  
Dev F1 tensor(0.3226)  
epoch 279, Loss: 0.636305800528563  
Dev F1 tensor(0.3226)  
epoch 280, Loss: 0.6361971032152176  
Dev F1 tensor(0.3226)  
epoch 281, Loss: 0.636090496787878  
Dev F1 tensor(0.3226)  
epoch 282, Loss: 0.63598419049183502  
Dev F1 tensor(0.3226)  
epoch 283, Loss: 0.6358793693734314  
Dev F1 tensor(0.3226)  
epoch 284, Loss: 0.6357759706909985  
Dev F1 tensor(0.3226)  
epoch 285, Loss: 0.63567368756870844  
Dev F1 tensor(0.3226)  
epoch 286, Loss: 0.6355726043987274  
Dev F1 tensor(0.3226)  
epoch 287, Loss: 0.63547300658512115  
Dev F1 tensor(0.3226)  
epoch 288, Loss: 0.63537581473920013  
Dev F1 tensor(0.3226)  
epoch 289, Loss: 0.635279633937033  
Dev F1 tensor(0.3226)  
epoch 290, Loss: 0.63518458943176  
Dev F1 tensor(0.3226)  
epoch 291, Loss: 0.63509039595599  
Dev F1 tensor(0.3226)  
epoch 29



Logistic Regression Results:

```
Accuracy: 0.84
F1-score tensor([0.8621])
```

## Written Assignment (60 Points)

Written assignment tests the understanding of the student for the assignment's task. We have split the writing into sections

will need to write 1-2 paragraphs describing the sections. Please be concise.

**In your own words, describe what the task is (20 points)**

Describe the task, how it is useful and an example.

**Section 1: Sentiment Classification Dataset**

We have two datasets of a collection of hotel reviews (positive and negative), not tokenized, or sentence segmented, and the words are space separated. V

(positive and negative), create training and test datasets (split the original dataset), and define functions for evaluation metrics (compute accuracy, precision, recall, and f1 score)

## Section 2: Baselines:

We need to create two baselines for Random Chance and Naïve Bayes Classifier. The Random Chance predict the label according to the labels' distribution. The Naïve Bayes Classifier using tokens in the training samples. We need preprocess the set and implement fit and predict methods. We will compare these two baselines by fitting on some sample set.

### Section 3: Logistic Regression on Features:

We need to build a logistic regression based on hand-engineered features. We need to implement feature extraction to transform the input data into a vector representation suitable for a classification system for sentiment analysis. We need to have a normalization formula to normalize dataset for comparison across different datasets. We need to implement weights and logistic loss function and convert logits into probabilities. We need to use different epochs and learning rate to improve the result.

**Describe your method for the task (10 points)**

Section 1: Sent

We use `util.py` to load the dataset. We use a for-loop to iterate items in labels: whenever we get `label == 0`, we increase `neg_pos` count by one, same logic for positive reviews. I use `pie chart` for visualization. For split the dataset, I transform the list object DataFrame first, then shuffle indices using `numpy.random.shuffle` to split the dataset randomly. I use `DataFrame.values.tolist()` to transform DataFrame back to list object.

Section 2: Baselines:

I use NumPy library to randomly generate value set based on training distributions. I use spacy to tokenize the text and help features from text. I preprocess text using lemma\_ function and remove stopwords using is\_stop function.

The fit method for Naïve Bayes Classifier I need to group samples by their labels and preprocess each text. I count the word for each label through a for-loop and dictionary, whenever we get a word, we update the dictionary. The predict method for Naïve Bayes Classifier I need to compute each word probability by calculate the count of that word divided by whole word of word. I use LaPlace smoothing, whenever I get a new word not in the training dictionary, I add 1 to it. I use math.log to sum all prob

### Section 3: Logistic Regression on Fe

I implement six features exactly showed in the table: count (positive/negative lexicon), count (1st and 2nd pronouns), add 1 'no' or '!', and log(word count of doc). Normalize torch.Tensor by iterating a 2D array and update values based on the formula into the weight using built-in library in torch, which is torch.nn.init.xavier\_uniform. I use built-in library in torch to implement logistic loss function, which is torch.nn.BCEWithLogitsLoss. I convert logits into prediction by using built-in library torch, which is torch.round.

### Experiment Results (10 points)

Typically a table summarizing all the different experiments

Section 1: Sentiment Classification Dataset:

All dataset

# of examples have label = 0	# of examples have label = 1
94	95

Train Label Distribution

	74	77
Dev Label Distribution		
	# of examples have label = 0	# of examples have label = 1
	20	18

Section 2: Baselines:

Dev Set Result

	Random Chance F1	Naïve Bayes F1
	0.5945945945945946	0.9090909090909091

Test Set Result (Deadline)

	Random Chance F1	Naïve Bayes F1
	0.491228	0.90196

Section 3: Logistic Regression on Features:  
With Epoch = 500 and learning rate = 0.5

Dev Set Result at Epoch 500		
	Loss	F1
	0.60736	0.9412

	Acc
	0.86842

Test Set Result/Logistic Regression Results (Deadline)

	Accuracy	F1-score tensor
	0.84	0.8621

Discussion (20 points)

Section 1: Sentiment Classification Dataset:

Use pie chart to visualize the distribution of positive and negative reviews. Although there are many other kinds of charts to the distribution, pie chart is the most intuitive ways for distribution. Use pandas library to split the dataset into training and by random shuffle the indices. However, there are many other ways to split the dataset but I'm more experience with pandas to transform original list object into DataFrame object than try other methods. Implement the evaluation metrics based on helps me to further understand accuracy, precision, recall and f1 score. I believe there are some libraries can automatically calculate these values.

Section 2: Baselines:

Use `numpy.random.choice` to generate 0 and 1's based on

tokenize sentences into words. Dictionary is quite useful when I implement fit method in Naive Bayes Classifier. I update dictionary by counting whenever I iterate a word in the sentence. The predict method I need to create another two dictionaries for probability of each word from previous counting dictionary. The Laplace smoothing method helps us to calculate probability when the word is not in our dictionary. I use log function because sum all log is much faster than multiple all the probabilities.

**Section 3: Logistic Regression on Features:**

Use three built-in library in torch.tensor to initialize weight combination of epoch and learning rate to get relatively weights and there are many more other built-in function can implement more engineering features to improve fe

can implement more engineering features to improve feature extraction, for example, the count of n-gram phonemes in the whole text, and the count of comma in the whole text, and etc. The combination of epochs and learning rate can be tuned using exist library parameter tuning.